


Bio-Inspired Artificial Intelligence

Lecture 6: Evolutionary Computing



<http://www.informatik.uni-hamburg.de/WTM/>

Motivation

- Problem solving is a central theme in computer science
 - Many new and more **complex problems**
 - Time for analysis and algorithms development decreases
 - More **universal** algorithms with automatic adaptation needed
 - “good” solutions within acceptable time are often satisfying
- Most powerful problem solvers in nature:
 - The (human) brain
 - ... that created “the wheel, New York, wars and so on” [Adams 1978]
 - The evolutionary mechanism
 - ... that created the human brain [Darwin 1895]

Computational Approach: Evolutionary Algorithms

Motivating Example



Goal of this lecture

- Learn about an approach to **approximate** problem solving based on the theory of evolution
- Get a deeper understanding of Evolutionary Algorithms





- Slides are mainly based on:
 - *Introduction to Evolutionary Computing* by Eiben and Smith, Springer 2003.
 - *Bio-Inspired Artificial Intelligence: Theories, Methods, and Technologies* by Floreano & Mattiussi, MIT 2008

The 4 pillars of Evolution

All species derive from
common ancestor

Charles Darwin, 1859
On the Origins of Species



- Population
 - Group of several individuals
- Diversity
 - Individuals have different characteristics
- Heredity
 - Characteristics are transmitted over generations
- Selection
 - Individuals produce more offspring than the environment can support
 - Better at food gathering = better at surviving = make more offspring

Evolutionary Computing

- Trial-and-error problem solving method
- Population of individuals with reproduction and mutation
- “Survival of the fittest” (Darwin)
- “Diversity drives change” (Darwin)
- Inspired by natural evolution:

Evolution		Problem Solving
Environment	↔	Problem
Individual	↔	Candidate solution
Fitness	↔	Quality


History

1948	Turing proposes “genetical or evolutionary search”
1962	Bremermann runs first computer experiments on “optimization through evolution and recombination”
1964	Rechenberg and Schwefel introduce <i>evolution strategies (ES)</i>
1965	Fogel, Owens and Walsh introduce <i>evolutionary programming (EP)</i>
1975	Holland introduces <i>genetic algorithms (GA)</i>
1992	Koza introduces <i>genetic programming (GP)</i>
1997	Launch of European EC Research Network EvoNet

- All techniques use same technology & ideas but differ in details
- Research area *Evolutionary Computing* deals with *Evolutionary Algorithms*

Biological Background

- Genotype determines Phenotype
- Genes: **complex mapping**
 - One gene may affect many traits (features)
 - Many genes may affect one trait
 - Small changes in the genotype lead to large changes in the organism (e.g. Height, hair colour)
 - Encoded in strands of DNA
- Human DNA is organised into chromosomes
 - Characterised by Alleles (allele is one of two or more forms of a gene or a gene locus)
 - Together define the physical attributes of the individual



Phenotype & Genotype


- **Phenotype**: Manifestation of the organism (appearance, behavior, etc.).
 - Selection operates on the phenotype;
 - It is affected by environment, development, and learning
- **Genotype**: The genetic material of that organism.
 - It is transmitted during reproduction;
 - It is affected by mutations;
 - Selection does not operate directly on it
- **Genetics**: Structure and operation of genes
- **Functional genomics**: Role of genes in the organism
- To what extent are we determined by genotype and phenotype?




Jean-Felix & Auguste Piccard

DNA (*DesoxyriboNucleic Acid*)

- Long molecule, twisted in a spiral and compressed
- Humans have 23 pairs of DNA molecules (*chromosomes*)
- DNA is composed of 2 complementary sequences (*strands*) of 4 nucleobases (A, T, C, G), which bind together in pairs (A-T and C-G)





- A **gene** is a sequence of several nucleotides that produce a protein




Cell Replication

- Cells replicate in two ways:
 - Mitosis**: during growth/maintenance of the organism
 - Meiosis**: during production of gametes (egg / sperm cells)



Cell Replication (cont.)

- Gametes: Sperm or egg cell
 - Contain only one single chromosome complement of chromosomes
 - Formed by a special form of cell splitting: **Meiosis**
 - During meiosis, pairs of chromosomes undergo **crossing-over**



- Occasionally some of the genetic material changes very slightly (replication error): **Mutations**

Example: Evolution of Camouflage



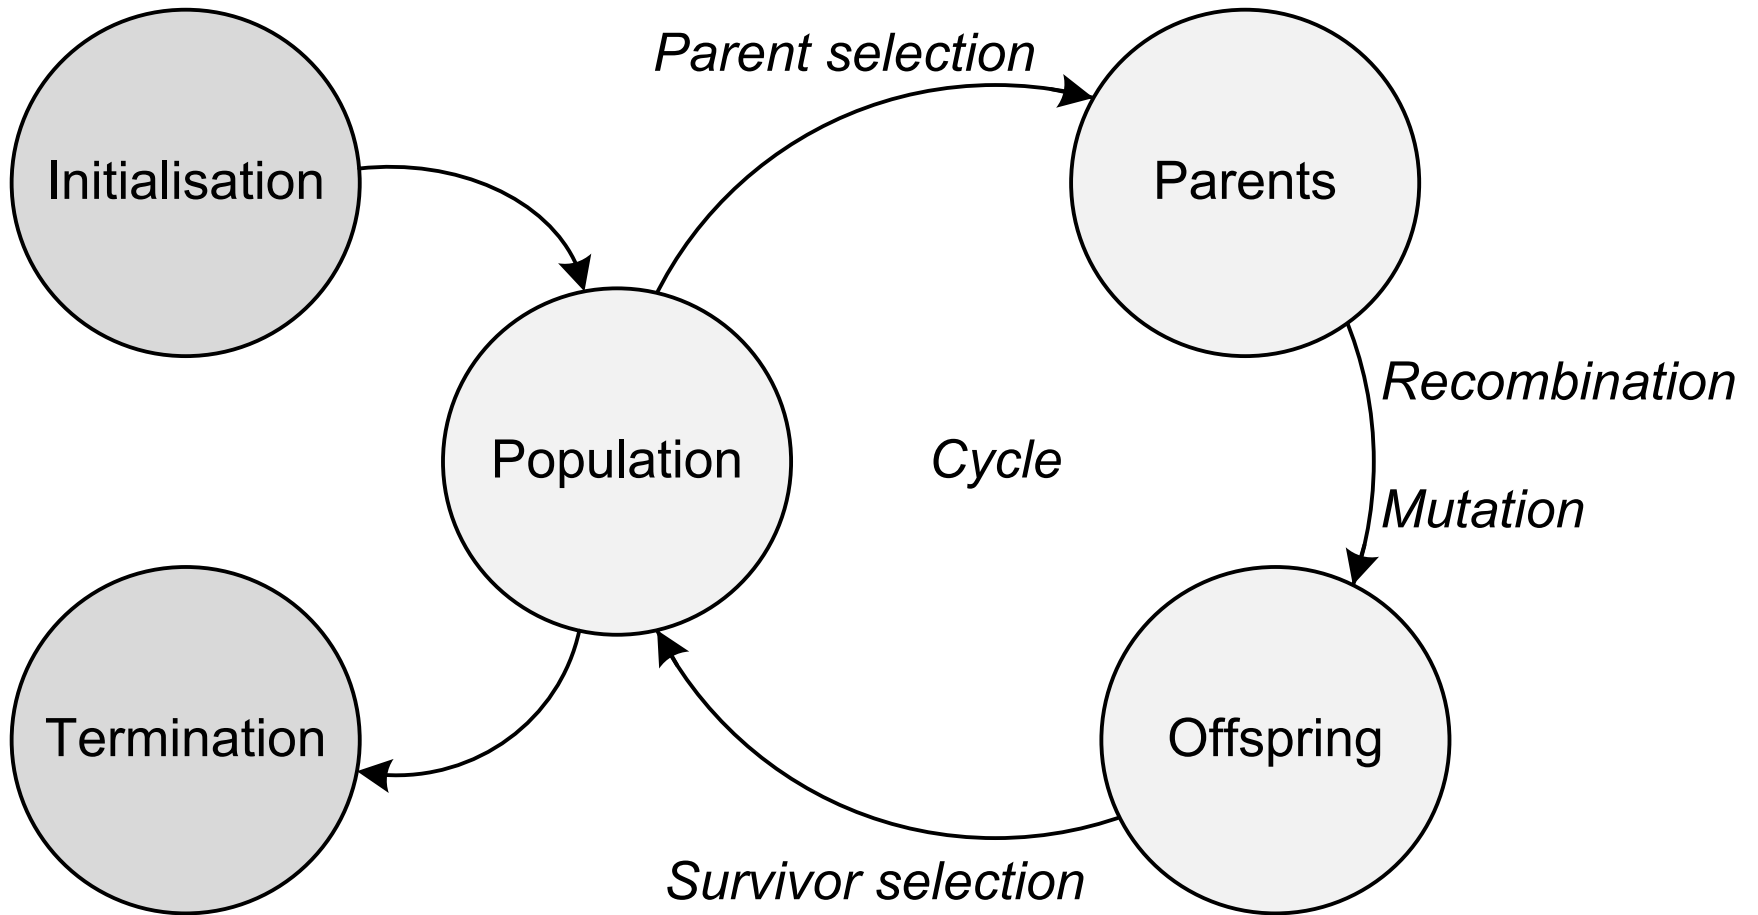
Chris Schneider, Boston University

General Characteristics of Evolutionary Algorithms (EA)

Evolutionary Algorithms are:

- generate-and-test algorithms
- population-based
- stochastic
- inspired by natural evolution
- search procedures:
 - guided by fitness
 - driven by variation and selection operators
- anytime algorithms: stop at any time with a (suboptimal) solution

General Scheme of EAs



General Categories of EAs

- Variation operators (recombination & mutation) depend on candidate representation
- Selection operators only require fitness value
 - “universal” operators
- EAs follow same basic scheme
 - differ mainly in technical details, e.g. representation of candidate solutions:

Algorithm Class	Candidate Representation
Genetic Algorithm	string over finite alphabet, e.g. 01110
Evolution Strategies	real-valued vectors, e.g. (0.1,0.4,12.3)
Evolutionary Programming	application specific, e.g. FSM
Genetic Programming	trees

Representation

- Cover all possible solutions
- Encoding should only allow valid solutions (not always possible)
- Choose ***appropriate representation***:
 - Bit-string can decode integers or real numbers
 - Problems: e.g. mutation changes values significantly (work-around: *Gray coding*: two successive values differ by 1 bit)
 - Better direct representation of numbers
 - Bit-string often fine
- Two meanings of representation:
 - Mapping between phenotype & genotype space: de-/encoding
 - Data structure used in genotype space

Fitness Function


- Represents requirements to adapt to
- Basis for selection
- Assigns quality measure to genotypes
- Synonyms: evaluation or objective function

Example:

- **Context**: Minimize x^2
- **Phenotype**: $x \in \mathbb{N}$
- **Genotype**: z : binary representation of x
- **Fitness Function**: fitness $f(z)$ of genotype z is defined as 1 divided by square of its corresponding phenotype, e.g.
 $z = 0010 \rightarrow \text{phenotype: } x = 2 \rightarrow f(z) = 1/x^2 = 0.25$

Fitness Landscape Metaphor

- Usually of a very *high dimension*
- Local optimum: solution better than neighbouring solutions
- Global optimum: best solution in landscape
- Uni-modal problem: only one local optimum
- Multi-modal problem: several local optima



Parent Selection

- Select individuals to create offspring
- Population level
- Often *probabilistic*:
 - Individuals with high fitness more likely to become parents
 - “Weak” individuals might also become parents (with low probability) to avoid local optima
 - Sum over all probabilities is 1.0
- Parent selection in ES: often chosen randomly from parent population
- Parent selection supports process of evolving better solutions over time

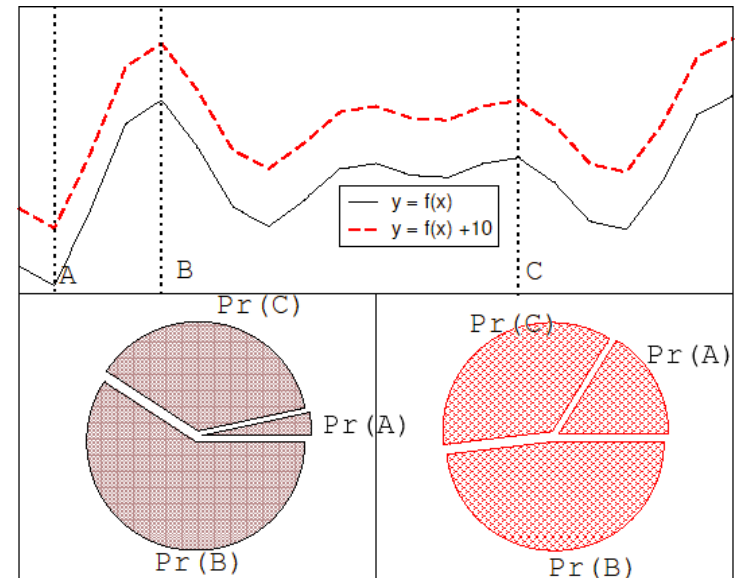
Parent Selection – Fitness Proportional Selection (FPS)

- Absolute fitness of individual vs. absolute fitness of population
- Probability to select i from population of size

$$\mu : \Pr(i) = f_i / \sum_i^{\mu} f_i$$

- Problems:

- **Premature convergence**
- Nearly equal fitness values result in no selection pressure
- Selection probabilities change for transposed fitness values
- Fitness scaling, e.g. $f(i) - \beta^t$ with $\beta = \min_{y \in P^t} f(y)$ (subtract minimum fitness in current population)



Parent Selection – Ranking Selection

- Inspired by problems of Fitness Proportional Selection
- Rank individuals by fitness
- Assign selection probabilities based on rank
- Different implementations of $Pr(i)$
- Problem: Can lead to slower convergence


Example: Linear Rank (LR)

$$P_{LR}(i) = \frac{(2-s)}{\mu} + \frac{2i(s-1)}{\mu(\mu-1)}$$

	Fitness	Rank	P_{selFP}	$P_{selLR}(s=2)$	$P_{selLR}(s=1.5)$
A	1	1	0.1	0	0.167
B	5	3	0.5	0.67	0.5
C	4	2	0.4	0.33	0.33
Sum	10		1.0	1.0	1.0

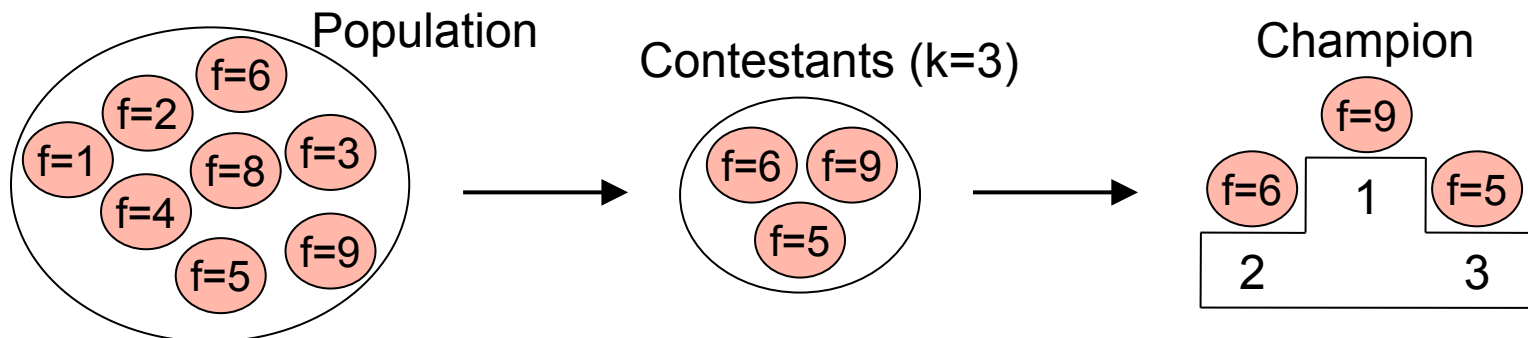
Parent Selection – Roulette Wheel

- FPS and Ranking Selection define probability distributions for selecting individuals
- How to sample these distributions?
 - *Roulette Wheel (RW)!*
- Sampling quality of RW in practice:
 - **High variance** from theoretical distribution
 - **Expensive** to calculate in
 - distributed systems
 - if number of parents μ is large



Parent Selection – Tournament Selection

- Tournament Selection only requires order between individuals (relative fitness) and no global knowledge
 - Previous methods use knowledge over entire population
- Idea:
 - randomly select k individuals into a group G of contestants
 - individual i with $f(i) = \max_{i \in G} f(i)$ wins the tournament
 - add i to parent set
 - repeat until μ parents are selected



k is the tournament size (larger size = larger selection pressure)


Variation: Recombination

- Variation operator: create new individual(s) from parents
 - Synonym: **crossover** operator
 - Distinguishes EA from other optimization techniques
 - Merge information from parents into offspring
 - Aims for diversity,
but sometimes a destructive jump in fitness landscape
- Crossover applied with probability p_c , e.g. $p_c \in [0.5, 1.0]$; otherwise parents are copied
- Implementation depends on representation form

Variation: Recombination – n -Point Crossover

- Split parents at n points and recombine segments
- **Positional** bias:
 - n -Point Crossover tends to keep together genes located close to each other
 - One-Point can never keep together genes from opposite ends
 - Knowledge on problem structure often not available


Example One-Point Crossover



Variation: Recombination – Uniform Crossover

- Swap genes based on vector of random values
 - Process generates first child
 - Second child: inverse process
- **Distributional** bias:
 - Genes are distributed among children instead of transferring larger sets of co-adapted genes to one child

Example: $p=0.5$, $x = (0.3, 0.7, 0.4, 0.2, 0.6, 0.9, 0.1, 0.8)$




Variation: Recombination – Further operators

■ Arithmetic Recombination:

- Powerful for floating-point representations


Example: Simple arithmetic recombination



■ Permutation:

- Powerful if exchange of substrings is a constraint

Example: Partially Mapped Crossover (PMX)




Variation: Mutation

- Variation operator: create new individual(s) from old ones
- Slightly mutates one individual
- Always *stochastic*: random and unbiased changes
- Implementation depends on representation form
- Many different implementations
- Like recombination, mutation plays different roles in different EAs

Variation: Mutation – Change of Allele Values




■ Bitwise mutation

- For every position: flip bit with probability p_m



■ Random resetting / Uniform mutation


- For every position: change value to random value from the corresponding domain with probability p_m



Variation: Mutation – Permutation


■ Swap mutation

- Randomly choose two positions and swap their allele values




■ Insert mutation

- Shuffle value of a second position towards the first position



■ Scramble mutation

- For a chosen subset scramble their positions




Survivor Selection

- Similarly to parent selection, survivor selection selects individuals based on quality
- Role: reduce μ parents and λ offspring to μ individuals that constitute the next generation
 - Fitter individuals more likely to survive
 - Weak individuals may also survive
- Selection often **deterministic** based on age and/or fitness
 - Age-based: Choose μ best from *offspring* only
 - Fitness-based: Choose μ best from *parents* and *offspring*
 - Elitism: Keep the $\kappa < \mu$ best, replace the rest by offspring
- Synonyms: environmental selection, replacement


Initialization

- Random initialization
- Heuristics:
 - might be used for higher quality initialization
 - additional time for implementation & computational effort
- EA should reach level of heuristic-based initialization after some k generations



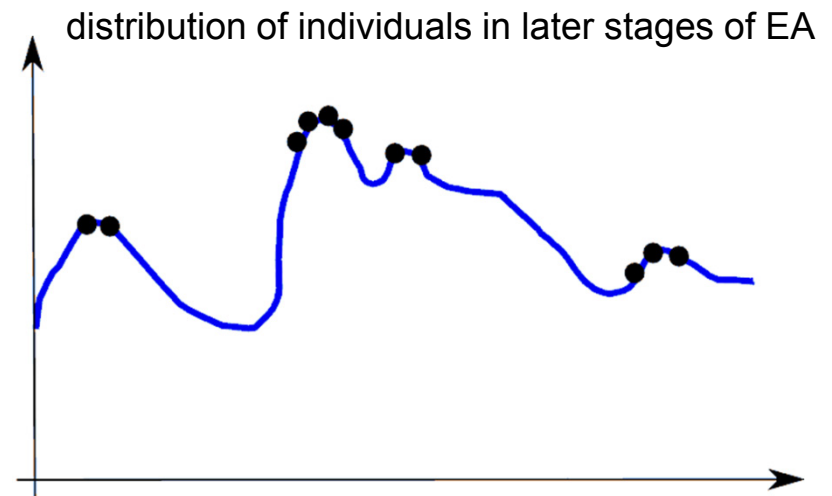

Termination Condition

- Optimum value known: stop if reached, or if only a small $\varepsilon > 0$ away
- Problem: EAs are stochastic and may never fulfill that condition
- Other possible criteria include:
 - CPU time elapsed
 - Maximum number of iterations or fitness calculations reached
 - No fitness improvement within last t generations
 - Diversity of population too small




Exploration vs. Exploitation

- **Exploration**: generate individuals in unexplored regions of the solution space
- **Exploitation**: generate individuals in regions with high fitness
- Trade-off:
 - too much exploration: inefficient search
 - too much exploitation: risk of getting stuck in local optima



Hybrid Evolutionary Algorithms

- Issue: For some problems an EA converges very slowly to a good solution
 - But: Some knowledge about solutions is available
 - Idea: Looking to improve EA search for good solutions
 - Hybridise EAs with **Local Search** operators that work within the EA loop
- Memetic Algorithm




Applications


- Civil, mechanical, and industrial engineering
- Optimal cutting
- Power systems
- Control systems
- Signal processing
- Rule for artificial player in computer games
- Machine learning
- ... and many more

Example: Travelling Salesman Problem

TSP30 Solution (Performance = 420)





TSP30 - Overview of Performance



Interesting program:
<http://www.lalena.com/AI/Tsp/>

Example: Antenna for Nanosatellites

- ST5 mission: Measure effect of solar activity on the Earth's magnetosphere
 - 3 nanosatellites (50 cm)
 - Design of antenna to send data to ground station



[Lohn, Hornby, Linden, 2004]

Genetic Representation

- Tree-based Encoding
- Execute instructions from root to leaves
- Evaluate fitness according to specs in simulation
- Build best and test in anechoic chamber

Function Set

f = forward (length, radius)



r x/y/z = rotate x/y/z

Terminal Set

Length, radius, x, y, z

Constraint:

max. 3 branches for each f node




Comparison human/evolved


Human




Evolved



AUT RHCP Theta Gain Bounds



Max and Min Gain vs Theta for 7.2 GHz



Example: Evolving of Vehicle Structures

BoxCar 2D

<http://boxcar2d.com/>

[Home](#) | [Designer](#) | [Best Cars](#) | [Forum](#) | [News](#) | [FAQ](#) | [The Algorithm](#) | [Versions](#) | [Contact](#)

Computation Intelligence Car Evolution Using Box2D Physics (v3.2)

60 fps average
Physics step: 2 ms (441 fps)
8 MB used

103

Generation: 3 Max Score: 103.9

Car	Score	Time
0	51.9	0:14
1	30.5	0:06
2	28	0:04
3	0.7	0:01
4	9.5	0:03
5	0.2	0:00
6	0.7	0:02
7	103.9	0:19
8	1.5	0:00
9	2.6	0:02
10	0	0:00
11	0	0:01
12	0.4	0:04
13	23.1	0:09
14	25.8	0:06
15	8.7	0:03
16	1	0:02
17	0.3	0:00
18	0.1	0:01



Time: 3:53

Score: 64.4

Torque: 71

max wheels wheel freq.

mutation rate

[Design a Car](#)

[See Changes / Play Older Versions](#)


Discussion

- Pros:
 - General heuristic
 - Able to find good solutions in feasible computing time
 - Distributed execution possible
- Cons:
 - Cannot guarantee optimal solutions
 - Long runtimes
 - Success depends also on used parameters

Further Readings:

- Eiben, A.E., Smith, J.E., *Introduction to Evolutionary Computing*, Springer, 2003

Outlook: Evolving of Creatures



Motivating Example – new Perspective

